

Worksheet -02

Outputs Screenshots

```
print("Dataset Preview:")
print(df.head())
print("\nDataset Information:")
print(df.info())
```

Dataset Preview:

	label	pixel_0	pixel_1	pixel_2	pixel_3	pixel_4	pixel_5	pixel_6	\
0	5	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	4	0	0	0	0	0	0	0	
3	1	0	0	0	0	0	0	0	
4	9	0	0	0	0	0	0	0	

	pixel_7	pixel_8	...	pixel_774	pixel_775	pixel_776	pixel_777	\
0	0	0	...	0	0	0	0	
1	0	0	...	0	0	0	0	
2	0	0	...	0	0	0	0	
3	0	0	...	0	0	0	0	
4	0	0	...	0	0	0	0	

	pixel_778	pixel_779	pixel_780	pixel_781	pixel_782	pixel_783
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

[5 rows x 785 columns]

Dataset Information:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60000 entries, 0 to 59999
Columns: 785 entries, label to pixel_783
dtypes: int64(785)
memory usage: 359.3 MB
None
```

```
▶ X = df.iloc[:, 1:-1].values
  y = df.iloc[:, -1].values

  label_encoder = LabelEncoder()
  y_encoded = label_encoder.fit_transform(y)

  one_hot_encoder = OneHotEncoder(sparse_output=False)
  y_one_hot = one_hot_encoder.fit_transform(y_encoded.reshape(-1, 1))

  print("\nUnique Classes:", np.unique(y))
  print("Encoded Labels:", np.unique(y_encoded))
  print("One-Hot Encoded Labels:\n", y_one_hot[:5])
```



```
Unique Classes: [0]
Encoded Labels: [0]
One-Hot Encoded Labels:
[[1.]
 [1.]
 [1.]
 [1.]
 [1.]]
```

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X, y_one_hot, test_size=0.2, random_state=

print("\nShapes:")
print("X_train:", X_train.shape, "y_train:", y_train.shape)
print("X_test:", X_test.shape, "y_test:", y_test.shape)
```



Shapes:
X_train: (48000, 783) y_train: (48000, 1)
X_test: (12000, 783) y_test: (12000, 1)

▼ Softmax Function:



```
import numpy as np

def softmax(z):
    """
    Compute the softmax probabilities for a given input matrix.

    Parameters:
    z (numpy.ndarray): Logits (raw scores) of shape (m, n), where
        - m is the number of samples.
        - n is the number of classes.

    Returns:
    numpy.ndarray: Softmax probability matrix of shape (m, n), where
        each row sums to 1 and represents the probability
        distribution over classes.

    Notes:
    - The input to softmax is typically computed as:  $z = XW + b$ .
    - Uses numerical stabilization by subtracting the max value per row.
    """
    z_shifted = z - np.max(z, axis=1, keepdims=True)
    exp_z = np.exp(z_shifted)
    return exp_z / np.sum(exp_z, axis=1, keepdims=True)
```

✓ Softmax Test Case:

This test case checks that each row in the resulting softmax probabilities sums to 1, which is the fundamental property of softmax.

```
[ ] z_test = np.array([[2.0, 1.0, 0.1], [1.0, 1.0, 1.0]])
    softmax_output = softmax(z_test)

    row_sums = np.sum(softmax_output, axis=1)

    assert np.allclose(row_sums, 1), f"Test failed: Row sums are {row_sums}"

    print("Softmax function passed the test case!")
```

➡ Softmax function passed the test case!

✓ Prediction Function:

```
[ ] def predict_softmax(X, W, b):
    """
    Predict the class labels for a set of samples using the trained softmax model.

    Parameters:
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d
    W (numpy.ndarray): Weight matrix of shape (d, c), where c is the number of classes.
    b (numpy.ndarray): Bias vector of shape (c,).

    Returns:
    numpy.ndarray: Predicted class labels of shape (n,), where each value is the index of the
    """

    z = np.dot(X, W) + b
    y_pred = softmax(z)

    predicted_classes = np.argmax(y_pred, axis=1)

    return predicted_classes
```

✓ Test Function for Prediction Function:

The test function ensures that the predicted class labels have the same number of elements as the input samples, verifying that the model produces a valid output shape.

```
▶ X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]])
  W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]])
  b_test = np.array([0.1, 0.2, 0.3])

  y_pred_test = predict_softmax(X_test, W_test, b_test)

  assert y_pred_test.shape == (3,), f"Test failed: Expected shape (3,), got {y_pred_test.shape}"

  print("Predicted class labels:", y_pred_test)
```

↗ Predicted class labels: [1 1 0]

✓ Loss Function:

```
[ ] def loss_softmax(y_pred, y):
    """
    Compute the cross-entropy loss for a single sample.

    Parameters:
    y_pred (numpy.ndarray): Predicted probabilities of shape (c,) for a single sample,
                           where c is the number of classes.
    y (numpy.ndarray): True labels (one-hot encoded) of shape (c,), where c is the number of c

    Returns:
    float: Cross-entropy loss for the given sample.
    """

    epsilon = 1e-12
    y_pred = np.clip(y_pred, epsilon, 1.0 - epsilon)
    n = y.shape[0]
    loss = -np.sum(y * np.log(y_pred)) / n
    return loss
```

✓ Test case for Loss Function:

This test case Compares loss for correct vs. incorrect predictions.

- Expects low loss for correct predictions.
- Expects high loss for incorrect predictions.

```
import numpy as np

y_true_correct = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]])
y_pred_correct = np.array([[0.9, 0.05, 0.05],
                           [0.1, 0.85, 0.05],
                           [0.05, 0.1, 0.85]])

y_pred_incorrect = np.array([[0.05, 0.05, 0.9],
                             [0.1, 0.05, 0.85],
                             [0.85, 0.1, 0.05]])

loss_correct = loss_softmax(y_pred_correct, y_true_correct)
loss_incorrect = loss_softmax(y_pred_incorrect, y_true_correct)

assert loss_correct < loss_incorrect, f"Test failed: Expected loss_correct < loss_incorrect, but got {loss_correct} > {loss_incorrect}"

print(f"Cross-Entropy Loss (Correct Predictions): {loss_correct:.4f}")
print(f"Cross-Entropy Loss (Incorrect Predictions): {loss_incorrect:.4f}")
```

```
→ Cross-Entropy Loss (Correct Predictions): 0.1435
Cross-Entropy Loss (Incorrect Predictions): 2.9957
```

✓ Cost Function:

```
[ ] def cost_softmax(X, y, W, b):  
    """  
    Compute the average softmax regression cost (cross-entropy loss) over all samples.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d), where n is the number of samples and d  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c), where n is the number of  
    W (numpy.ndarray): Weight matrix of shape (d, c).  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    float: Average softmax cost (cross-entropy loss) over all samples.  
    """  
  
    n = X.shape[0]  
    z = np.dot(X, W) + b  
    y_pred = softmax(z)  
    cost = loss_softmax(y_pred, y)  
    return cost
```

✓ Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

✓ Test Case for Cost Function:

The test case assures that the cost for the incorrect prediction should be higher than for the correct prediction, confirming that the cost function behaves as expected.

```
import numpy as np

X_correct = np.array([[1.0, 0.0], [0.0, 1.0]])
y_correct = np.array([[1, 0], [0, 1]])
W_correct = np.array([[5.0, -2.0], [-3.0, 5.0]])
b_correct = np.array([0.1, 0.1])

X_incorrect = np.array([[0.1, 0.9], [0.8, 0.2]])
y_incorrect = np.array([[1, 0], [0, 1]])
W_incorrect = np.array([[0.1, 2.0], [1.5, 0.3]])
b_incorrect = np.array([0.5, 0.6])

cost_correct = cost_softmax(X_correct, y_correct, W_correct, b_correct)

cost_incorrect = cost_softmax(X_incorrect, y_incorrect, W_incorrect, b_incorrect)

assert cost_incorrect > cost_correct, f"Test failed: Incorrect cost {cost_incorrect} is not gr

print("Cost for correct prediction:", cost_correct)
print("Cost for incorrect prediction:", cost_incorrect)

print("Test passed!")
```

```
→ Cost for correct prediction: 0.0006234364133349324
Cost for incorrect prediction: 0.29930861359446115
Test passed!
```


✓ Computing Gradients:

```
def compute_gradient_softmax(X, y, W, b):  
    """  
    Compute the gradients of the cost function with respect to weights and biases.  
  
    Parameters:  
    X (numpy.ndarray): Feature matrix of shape (n, d).  
    y (numpy.ndarray): True labels (one-hot encoded) of shape (n, c).  
    W (numpy.ndarray): Weight matrix of shape (d, c).  
    b (numpy.ndarray): Bias vector of shape (c,).  
  
    Returns:  
    tuple: Gradients with respect to weights (d, c) and biases (c,).  
    """  
  
    n, d = X.shape  
    z = np.dot(X, W) + b  
    y_pred = softmax(z)  
  
    grad_W = np.dot(X.T, (y_pred - y)) / n  
    grad_b = np.sum(y_pred - y, axis=0) / n  
  
    return grad_W, grad_b
```

The test checks if the gradients from the function are close enough to the manually computed gradients using `np.allclose`, which accounts for potential floating-point discrepancies.

```
[ ] import numpy as np

X_test = np.array([[0.2, 0.8], [0.5, 0.5], [0.9, 0.1]])
y_test = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]])

W_test = np.array([[0.4, 0.2, 0.1], [0.3, 0.7, 0.5]])
b_test = np.array([0.1, 0.2, 0.3])

grad_W, grad_b = compute_gradient_softmax(X_test, y_test, W_test, b_test)

z_test = np.dot(X_test, W_test) + b_test
y_pred_test = softmax(z_test)

grad_W_manual = np.dot(X_test.T, (y_pred_test - y_test)) / X_test.shape[0]
grad_b_manual = np.sum(y_pred_test - y_test, axis=0) / X_test.shape[0]

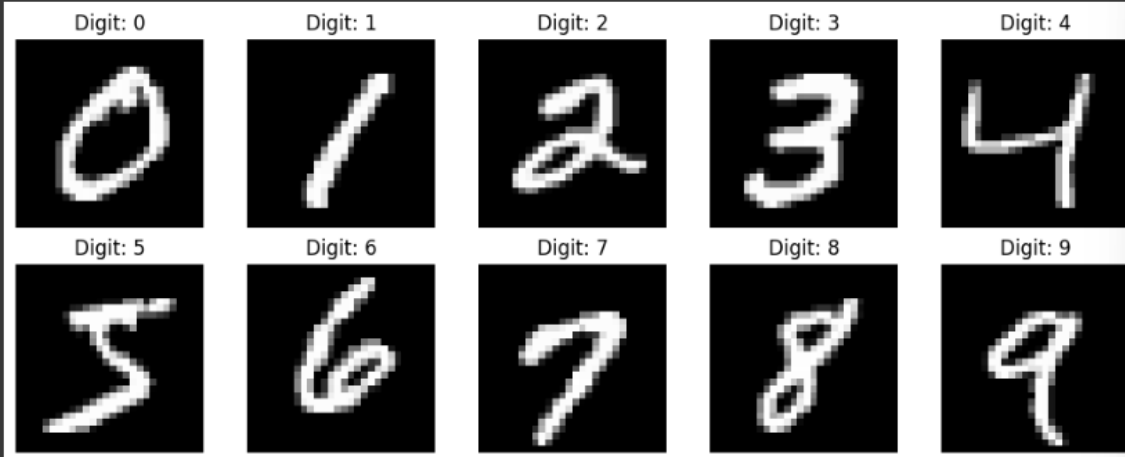
assert np.allclose(grad_W, grad_W_manual), f"Test failed: Gradients w.r.t. W are not equal.\nE
assert np.allclose(grad_b, grad_b_manual), f"Test failed: Gradients w.r.t. b are not equal.\nE

print("Gradient w.r.t. W:", grad_W)
print("Gradient w.r.t. b:", grad_b)

print("Test passed!")
```

```
↩ Gradient w.r.t. W: [[ 0.1031051  0.01805685 -0.12116196]
 [-0.13600547  0.00679023  0.12921524]]
Gradient w.r.t. b: [-0.03290036  0.02484708  0.00805328]
Test passed!
```

```
[ ] csv_file_path = "/content/drive/MyDrive/Sem6/AI and ML Workshop/Week-2/mnist_dataset.csv"
X_train, X_test, y_train, y_test = load_and_prepare_mnist(csv_file_path)
```



✓ A Quick debugging Step:

```
[ ] assert len(X_train) == len(y_train), f"Error: X and y have different lengths! X={len(X_train)}"
print("Move forward: Dimension of Feature Matrix X and label vector y matched.")
```



Move forward: Dimension of Feature Matrix X and label vector y matched.


▼ Train the Model:

```
[ ] print(f"Training data shape: {X_train.shape}")  
    print(f"Test data shape: {X_test.shape}")
```

↗ Training data shape: (48000, 784)
Test data shape: (12000, 784)

```
from sklearn.preprocessing import OneHotEncoder  
  
if len(y_train.shape) == 1:  
    encoder = OneHotEncoder(sparse_output=False)  
    y_train = encoder.fit_transform(y_train.reshape(-1, 1))  
    y_test = encoder.transform(y_test.reshape(-1, 1))  
  
d = X_train.shape[1]  
c = y_train.shape[1]  
  
W = np.random.randn(d, c) * 0.01  
b = np.zeros(c)  
  
alpha = 0.1  
n_iter = 1000  
  
W_opt, b_opt, cost_history = gradient_descent_softmax(X_train, y_train, W, b, alpha, n_iter, ε)  
  
plt.plot(cost_history)  
plt.title('Cost Function vs. Iterations')  
plt.xlabel('Iterations')  
plt.ylabel('Cost')  
plt.grid(True)
```

```
Iteration 0: Cost = 2.199553
Iteration 100: Cost = 0.607084
Iteration 200: Cost = 0.489327
Iteration 300: Cost = 0.440739
Iteration 400: Cost = 0.412682
Iteration 500: Cost = 0.393835
Iteration 600: Cost = 0.380030
Iteration 700: Cost = 0.369336
Iteration 800: Cost = 0.360723
Iteration 900: Cost = 0.353583
Iteration 999: Cost = 0.347588
```



```

[ ] ax.set_xticks(range(num_classes))
    ax.set_yticks(range(num_classes))
    ax.set_xticklabels([f'Predicted {i}' for i in range(num_classes)])
    ax.set_yticklabels([f'Actual {i}' for i in range(num_classes)])

    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            ax.text(j, i, cm[i, j], ha='center', va='center', color='white' if cm[i, j] > np.max(

    ax.grid(False)
    plt.title('Confusion Matrix', fontsize=14)
    plt.xlabel('Predicted Label', fontsize=12)
    plt.ylabel('Actual Label', fontsize=12)

    plt.tight_layout()
    plt.colorbar(cax)
    plt.show()

```



```


Confusion Matrix:
[[1127   0   5   2   3   11   9   2   13   3]
 [  0 1275   7  10   1   5   1   4  18   1]
 [  2  15 1026  16  20   3  27  24  33   8]
 [  8   6  33 1050   1  52   9   8  31  21]
 [  1   5   7   1 1092   0  10   4   4  52]
 [ 22  14  14  44  11  920  14   7  44  14]
 [  7   2   9   1  11  16 1119   2  10   0]
 [  6  27  26   4  15   3   0 1181   7  30]
 [  8  29  15  33   9  33  12   6 1000  15]
 [  8   6  10  18  44   9   0  38  10 1051]]

Precision: 0.90
Recall: 0.90
F1-Score: 0.90

```



Confusion Matrix



1200

1000

800

600

400

200

```
<ipython-input-24-6a5a1fc6a1b>:19: UserWarning: You passed a edgecolor/edgecolors ('k') for an  
plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test, cmap='coolwarm', edgecolors='k', marker='x')
```

[]

